

GRID SEARCH CV

1. The Client Requirement Input is provides the Data Set of Insurance and Output is he wants to predict the Insurance charges based on the several parameters.

2. The Basic Information is,

Input →Dataset, Output →Insurance charges

Total No of Rows= 1338

Total No of Columns= 6

3. Here the Preprocessing Method is, to handle Categorical column using, Converting String to Number (Nominal Data→ One Hot Encoder)

1. Support Vector MachineGRIDCV

```
from sklearn.model_selection import GridSearchCV
```

```
from sklearn.svm import SVR
```

```
param_grid = {'kernel':['rbf','poly','sigmoid','linear'],
```

```
               'C':[10,100,1000,2000,3000],'gamma':['auto','scale']}
```

```
grid = GridSearchCV(SVR(), param_grid, refit = True, verbose = 3,n_jobs=-1)
```

```
# fitting the model for grid search
```

```
grid.fit(X_train, y_train)
```

```
The R_score value for best parameter {'C': 3000, 'gamma': 'scale', 'kernel':  
'poly'}: 0.8577898390283539
```

3. Decision Tree

```
from sklearn.tree import DecisionTreeRegressor
```

```
from sklearn.model_selection import GridSearchCV
```

```

from sklearn.tree import DecisionTreeRegressor

param_grid = {'criterion':['mse','mae','friedman_mse','poisson'],

              'max_features': ['auto','sqrt','log2'],

              'splitter':['best','random']}

grid = GridSearchCV(DecisionTreeRegressor(), param_grid, refit = True, verbose = 3,n_jobs=-1)

# fitting the model for grid search

grid.fit(X_train,y_train)

The R_score value for best parameter {'criterion': 'friedman_mse', 'max_features': 'sqrt', 'splitter': 'best'}: 0.7310420392095164

```

4. Random Forest

```

from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestRegressor
param_grid = {'random_state': [0],
              'n_estimators':[100],
              'criterion':['mse','mae','friedman_mse','poisson']}
grid = GridSearchCV(RandomForestRegressor(), param_grid, refit = True, verbose = 3,n_jobs=-1)

# fitting the model for grid search
grid.fit(X_train, y_train)

The R_score value for best parameter {'criterion': 'friedman_mse', 'n_estimators': 100, 'random_state': 0}: 0.8595964959543749

```

5. Ada Boost Algorithm

```

from sklearn.model_selection import GridSearchCV

from sklearn.ensemble import AdaBoostRegressor

param_grid = {'random_state': [0],

              'n_estimators':[100], 'loss':['linear'],'learning_rate':[1.0]}

grid=GridSearchCV(AdaBoostRegressor(),param_grid,refit=True,verbose = 3,n_jobs=-1)

```

```
grid.fit(X_train,y_train)
```

The R_score value for best parameter {'learning_rate': 1.0, 'loss': 'linear', 'n_estimators': 100, 'random_state': 0}: **0.8618083410771146**

6.XGBoost Algorithm

```
from sklearn.model_selection import GridSearchCV
from xgboost import XGBRegressor
param_grid = {'n_estimators':[1000], 'max_depth':[7], 'eta':[0.1], 'subsample':[0.7], 'colsample_bytree':[0.8], 'use_rmm':['true'],
               'booster':['gbtree'], 'device':['cpu'],
               'verbosity':[1], 'validate_parameters':['false'], 'disable_default_eval_metric':['false'],
               'gamma':[0], 'min_child_weight':[1], 'max_delta_step':[0], 'sampling_method':['uniform'],
               'colsample_bylevel':[1], 'colsample_bynode':[1], 'Lambda':[1], 'alpha':[1], 'tree_method':
               ':[\"auto\"],
               'scale_pos_weight':[1], 'refresh_leaf':[1], 'process_type':['default'], 'grow_policy':['depthwise'], 'max_leaves':[0], 'max_bin':[256],
               'num_parallel_tree':[1], 'save_period':[0], 'task':['train'], 'model_in':['NULL'], 'model_out':['NULL'], 'model_dir':['models/'], 'dump_format':['text'],
               'name_dump':['dump.txt'], 'name_pred':['pred.txt'], 'pred_margin':[0]}
grid = GridSearchCV(XGBRegressor(), param_grid, refit = True, verbose = 3, n_jobs=-1)

grid.fit(X_train,y_train)
```

The R_score value for best parameter {'Lambda': 1, 'alpha': 1, 'booster': 'gbtree', 'colsample_bylevel': 1, 'colsample_bynode': 1, 'colsample_bytree': 0.8, 'device': 'cpu', 'disable_default_eval_metric': 'false', 'dump_format': 'text', 'eta': 0.1, 'gamma': 0, 'grow_policy': 'depthwise', 'max_bin': 256, 'max_delta_step': 0, 'max_depth': 7, 'max_leaves': 0, 'min_child_weight': 1, 'model_dir': 'models/', 'model_in': 'NULL', 'model_out': 'NULL', 'n_estimators': 1000, 'name_dump': 'dump.txt', 'name_pred': 'pred.txt', 'num_parallel_tree': 1, 'pred_margin': 0, 'process_type': 'default', 'refresh_leaf': 1, 'sampling_method': 'uniform', 'save_period': 0, 'scale_pos_weight': 1, 'subsample': 0.7, 'task': 'train', 'tree_method': 'auto', 'use_rmm': 'true', 'validate_parameters': 'false', 'verbosity': 1}: **0.832608353523225**

7.LG Boost Algorithm

```
param_grid = {
    'boosting_type': ['gbdt'],
    'num_leaves': [31],
    'max_depth': [-1],
    'learning_rate': [0.1],
    'n_estimators': [100],
    'subsample_for_bin': [200000],
    'objective': [None], # Assuming you meant to pass None, not 'None'
    'class_weight': [None],
    'min_split_gain': [0.0],
    'min_child_weight': [0.001],
    'min_child_samples': [20],
    'subsample': [1.0],
    'subsample_freq': [0],
    'force_row_wise': [True], # Use Python boolean True instead of "true"
    'colsample_bytree': [1.0],
    'reg_alpha': [0.0],
    'reg_lambda': [0.0],
    'random_state': [None],
    'n_jobs': [None], # Use None without quotes
    'importance_type': ['split']
}
```

```
grid = GridSearchCV(LGBMRegressor(), param_grid, refit = True, verbose = 3,n_jobs=-1)
grid.fit(X_train, y_train)
```

The R_score value for best parameter {'boosting_type': 'gbdt', 'class_weight': None, 'colsample_bytree': 1.0, 'force_row_wise': True, 'importance_type': 'split', 'learning_rate': 0.1, 'max_depth': -1, 'min_child_samples': 20, 'min_child_weight': 0.001, 'min_split_gain': 0.0, 'n_estimators': 100, 'n_jobs': None, 'num_leaves': 31, 'objective': None, 'random_state': None, 'reg_alpha': 0.0, 'reg_lambda': 0.0, 'subsample': 1.0, 'subsample_for_bin': 200000, 'subsample_freq': 0}: **0.8699321391117371**

The Result of LG Boost Algorithm is better accuracy value 0.869 compared to all Algorithm.

